

A Tiered Compilation Architecture for Low-Latency Wasm Runtimes in Edge Computing

Master's Thesis Defense

Hsiang-Yu Lee (李祥宇)

Advisor: Prof. Shih-Wei Liao

June 25, 2026

Department of Computer Science and Information Engineering
National Taiwan University

Outline

Motivation

Problem Statement and Contributions

Architecture Overview

Tier-1: Baseline JIT

Tier-2: Selective LLVM Recompilation

On-Stack Replacement

Evaluation

Limitations and Future Work

Conclusion

Motivation


Edge computing reshapes the workload

Edge nodes face two tight budgets.

- **Response time:** tens of milliseconds per request.
- **Resources:** 1–2 GB RAM and few vCPUs.

The workload mix is short-lived, stateless functions that start and exit constantly.

WebAssembly fits the role. The format is compact, language-agnostic, sandboxed, and predictable. Fastly Compute, Cloudflare Workers, and Fermion Spin all run WebAssembly in production.



Per-request lifecycle.

The WasmEdge cold-start bottleneck

WasmEdge is the leading server-side Wasm runtime. Microsoft, ByteDance, and Sony all run WasmEdge in production.

WasmEdge compiles with LLVM end-to-end. The LLVM pipeline produces near-native code. The pipeline is also expensive:

- Compile time grows faster than module size.
- Memory usage during compilation is large.
- Every cold start pays the full compile cost.

The structural problem. A single-tier compiler treats every invocation the same way. A request that runs for one microsecond pays the same compile cost as one that runs for one minute.

Definition

Time from request arrival to the first machine-code instruction of the requested wasm function. Assumes no ready-to-run compiled instance of the module is resident.

Inside the window

- load, parse, validate
- compile
- instantiate

Outside the window

- the wasm function's own execution
- reported as steady-state work

Execution time is highly skewed

A small fraction of functions accounts for the bulk of run-time.

- Heavy optimization on cold functions wastes the compile budget.
- Cold functions receive optimization they do not need.
- Hot functions wait behind cold compilation before they can start.

The established solution is tiered compilation.

- A **Tier-1** baseline compiles every function fast.
- A **Tier-2** optimizer recompiles hot functions later, guided by profile data.
- V8, JavaScriptCore, and SpiderMonkey all use this pattern.

WasmEdge has no baseline tier, no profile-driven promotion, and no mid-execution code replacement.

Problem Statement and Contributions

Three sub-problems

Goal. Cut cold-start latency without giving up the steady-state throughput that LLVM delivers on hot code.

Tier-1 baseline. Compile every wasm function fast at module load. Cover the full WebAssembly instruction set the workloads use. Reuse the existing WasmEdge runtime contracts.

Tier-2 optimizer. Reuse the LLVM frontend to recompile hot functions. Bridge the two tiers' calling conventions. Hide the LLVM compile cost from the user thread.

In-flight migration. A function called once spends its run-time inside one long loop. The prologue never re-fires, so entry-table promotion delivers no speedup. A separate mechanism must migrate the running frame mid-loop.

Contributions

1. **Tier-1 baseline JIT for WasmEdge.**

Integrate a lightweight Sea-of-Nodes library.

Translate wasm bytecode to sea-of-nodes IR nodes.

Embed per-function and per-loop profile counters directly in compiled code.

2. **Tier-2 selective LLVM-frontend recompilation.**

Design a heuristic to synthesize a mini-module per batch.

Design three ABI bridge thunks to bridge the calling convention gap between the two tiers.

3. **On-Stack Replacement.**

Instrument every outermost loop with a three-state back-edge diamond.


Synthesize a continuation that starts execution at the loop header.

4. **Empirical evaluation.**

Report and discuss a step-by-step experiment on 39 strengthened sightglass kernels.

Architecture Overview

System architecture



Pipeline at a glance


- The user thread runs **Tier-1** native code and trips the embedded counters.
- A heuristic selects a batch of functions to promote to **Tier-2** and synthesizes a mini-module for the batch.
- Two promotion paths feed two queues: function-entry promotion and loop OSR.
- A worker thread compiles **Tier-2** batches and OSR continuations in the background.
- The worker publishes each result with one atomic store into the dispatch slot.

Tier-1: Baseline JIT

What is Sea-of-Nodes?


A function becomes **one graph**: nodes are operations, edges are dependencies (data *and* control). Most nodes *float* until a late pass schedules them into blocks.

Data: $d = (a + b) * c$



Each value is a node; edges are data deps.

Control + data: if p : $x=1$ else 2



ϕ merges the two branch values — SSA built in.

Tier-1: a lightweight Sea-of-Nodes baseline

Code generator: dstogov/ir. A small Sea-of-Nodes JIT library, descended from the Zend JIT in PHP 8.

- The library runs a fixed, linear pipeline. No pass manager.
- Compile latency stays bounded and predictable.
- The codebase is small enough to vendor and modify.

Advantages of sea-of-nodes IR.

- **SSA by construction.** Values are graph nodes. Definitions and uses are explicit edges. No separate SSA-construction phase.
- **One graph for control flow and data.** A single traversal covers both. Passes do not need parallel analyses.
- **Cheap global code motion.** Floating nodes let the scheduler choose placement after optimization.

Tier-1: from wasm to native

Each wasm function is compiled at module instantiation, in three phases.

Phase 1 — one forward walk over the body (*interleaved, per instruction*):

- **Operand stack.** Model the wasm value stack as a stack of IR refs: pop operands, emit a node, push the result.
- **Locals** → **SSA.** Track each local's current node; `local.get` reads it, `local.set` updates it. PHIs only at loop headers, only for loop-modified locals.
- **Instrumentation.** Prologue call counter; three-state OSR diamond on each outermost-loop back-edge.

⇒ a complete Sea-of-Nodes SSA graph.

Phase 2 — dstogov/ir backend. Fold / value-number, constant propagation, global code motion, scheduling, register allocation, code emission.

Phase 3 — publish. Store the native entry pointer in `FuncTable[FuncIdx]`.

Tier-2: Selective LLVM Recompilation

Mini-module synthesis & Promotion

Tier-2 compiles a batch. A batch is a small set of hot or hot-adjacent functions selected from the original module. Functions outside the batch become stubs. Thunks later rewrite those stubs into cross-tier calls.

Algorithm 1

1. Deep-copy the original wasm module.
2. Replace every defined function outside the batch with a small stub.
3. Run the LLVM frontend on the mini-module unoptimized.
4. Post-process the stubs with thunks, and prune unused stubs.
5. Run the LLVM optimizer on the mini-module.
6. Emit native code and publish the `fwd_thunk` pointer into `FuncTable`.

We run the LLVM frontend unoptimized before post-processing to preserve the call sites that post-processing rewrites into thunk calls.

Three ABI bridge thunks

Tier-1 passes arguments in a flat 64-bit buffer. LLVM passes arguments in typed registers. Three thunks bridge the mismatch, one per direction of cross-tier dispatch.

Thunk	Reached via	Role
fwd_thunk	Tier-1 via FuncTable slot	Repacks the flat buffer into typed params, then calls the Tier-2 body.
t1_thunk	Tier-2 direct call (by name)	The rewritten non-batch stub; repacks typed args back into the Tier-1 buffer.
entry_thunk	LLVM indirect call (by pointer)	The Tier-1 function's table entry; catches call_indirect from Tier-2 /OSR/AOT.

Each thunk handles one direction. Neither tier needs to know the other tier's ABI.

Promoting only the threshold-tripping function is not enough. The function that trips the counter is often a small leaf inside a larger hot caller. A leaf promoted alone leaves its caller in **Tier-1**, so the optimizer cannot inline across the boundary.

Algorithm 2

1. **Walk up one hop.** Pick the hottest direct caller as the batch root.
2. **One-hop BFS.** Scan the root's direct callees. A batch size limit caps the count.
3. **Two-test admission** for each candidate callee:
 - Dynamic test: call counter > 0 — filters dead code (never called).
 - Static test: caller references the callee ≥ 2 — filters one-off helpers (init/cleanup).

A callee joins the batch if *either* test passes.

On-Stack Replacement

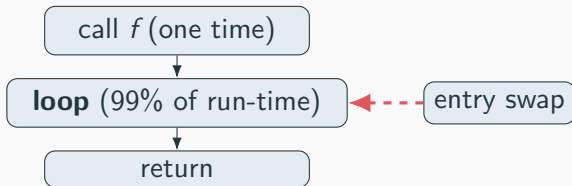
The one-shot-caller problem

A function called once, spending its run-time inside one long loop.

This is the shape of cryptographic kernels, parsing loops, and many serverless handlers.

Function-entry promotion delivers no speedup.


- Entry promotion redirects the next call to the function.
- This function is called once. There is no next call.
- The **Tier-1** frame keeps running the loop until return.



no future call to redirect

Three-state back-edge instrumentation

Each outermost loop reads one slot of `OsEntryTable` on every back-edge. The slot encodes a three-state machine.



The tier-2 promotion queue and the OSR promotion queue share one worker. The OSR queue is drained first.

Evaluation

Benchmark suite: sightglass-strong.

Upstream sightglass kernels run in well under a second. The strengthened suite extends each kernel's **Tier-1**-only run-time into a 5–10 s band. **Tier-2** now has enough post-switch runway to amortize the worker's compile cost.

39 analyzed kernels group into five execution shapes.

- **Frequent-call** (17): one hot function called many times.
- **Hot-leaf** (6): a small helper is called many times and becomes hot before the hot caller.
- **One-shot-caller** (7): outer function called once, dominated by one loop.
- **Resistant** (4): too little run-time or no inline opportunity.
- **Compile-dominated** (7): LLVM compile cost dwarfs the wasm work.

Three configurations, three steps


Each step turns on one mechanism on top of the previous step.

Step	t1/t2	LLVM/t2	Mechanism added
Step 1 – promote hot callees only	0.97×	0.76×	function-entry promotion
Step 2 – walk-up + bounded BFS	1.03×	0.80×	better batch composition
Step 3 – on-stack replacement (full system)	1.11×	0.86×	in-flight migration

- Geometric mean across all 39 kernels.
- t1/t2: how much faster the full system runs than **Tier-1** alone.
- LLVM/t2: how close the full system gets to whole-module LLVM JIT.
A value of 0.86× means the full system recovers 86% of LLVM throughput.

Step 2: walk-up + BFS


Step 1 batched only the leaf that tripped the counter. Step 2 rolls the surrounding caller into the batch.



Still lagging: one-shot-caller kernels (shootout-ctype 0.37 $\times$, blind-sig 0.40 $\times$) need OSR ^{21/35}

Step 3: OSR closes the one-shot-caller gap

The seven kernels with the largest Step 2 → Step 3 jump all sit in the one-shot-caller group.



At the swing: the counter saturates mid-loop, and the worker's continuation runs the rest optimized.

Compile time and end-to-end wall-clock

Compile time. How long the runtime spends compiling before steady state.

End-to-end wall-clock. How long the user actually waits.

Compile-time gain

Geometric mean: **11.1**×

The full system finishes compilation about 11 times faster than whole-module LLVM JIT.

The **Tier-2** optimizer runs on the worker thread, off the user's critical path.

End-to-end gain

Geometric mean: **1.73**×

The full system finishes the end-to-end workload about 73% faster on average.

The gain comes from saved compile time, even though steady-state throughput is 14% lower.

The largest wall-clock wins

The wall-clock benefit is largest on compile-dominated workloads. These workloads are the real-world targets for the architecture.

Kernel	Tier-2 time	LLVM time	speedup
tinygo-json	473 ms	10 361 ms	21.9×
tract-onnx-image-classification	9.7 s	136 s	14.1×
spidermonkey-json	~9 s	~90 s	~10×
spidermonkey-markdown	~9 s	~90 s	~10×
spidermonkey-regex	~9 s	~90 s	~10×
image-classification	303 ms	2 667 ms	8.8×
sqlite3	1 985 ms	16 953 ms	8.5×

Why the wins cluster here. Each module is interpreter-shaped, reflection-heavy, host-bound, or dominated by large tensor operators. The source-to-wasm toolchain has already baked in the hot path, so LLVM cannot recover further speedup. Saving the LLVM compile cost is therefore pure win.

Where the tiered design loses ground

Long-running synthetic kernels. The **Tier-1** throughput gap is no longer offset by a meaningful cold-start saving.

Kernel	end-to-end vs. LLVM
shootout-ctype	0.93×
shootout-sieve	0.84×
shootout-fib2	0.90×

The central observation. Workloads split bimodally. Compile-dominated workloads win big.

Long-running synthetic kernels lose a small amount. The architecture trades a modest steady-state cost for a large cold-start saving. The trade pays off most clearly on the real-world workloads the architecture targets.

Limitations and Future Work

No SIMD lowering. **Tier-1** does not lower wasm SIMD instructions. The translator rejects modules that use them. The **Tier-2** promotion gate accepts only scalar signatures.

No de-tiering. Promotion is one-way. A function promoted under one workload shape stays promoted even if the shape changes.

x86-64 only. The current prototype runs on x86-64. The library backend supports aarch64, but no aarch64 port has been validated.

Conclusion

Summary

This thesis presents a **tiered compilation architecture for WasmEdge** that separates cold-start latency from peak code quality.

Three design components.

- **Tier-1**: Sea-of-Nodes baseline JIT with embedded counters and a uniform calling convention.
- **Tier-2**: mini-module synthesis drives the LLVM recompilation; three ABI bridge thunks; walk-up batch policy.
- **OSR**: three-state back-edge diamond; continuation resumes a running frame mid-loop.

Headline results across 39 kernels.

- Compile time: **11.1**× faster than whole-module LLVM JIT.
- Steady-state throughput: **86%** of LLVM JIT preserved.
- End-to-end wall-clock: **1.73**× faster on average; up to 21.9× on `tinygo-json`.

Thank You

Questions?

Author: Hsiang-Yu Lee (李祥宇)

Advisor: Prof. Shih-Wei Liao

Code: WasmEdge fork with tiered JIT

Calling convention and JitExecEnv

Every **Tier-1** function takes the same two arguments.

$$f(\text{env}^*, \text{args}^*) \rightarrow \text{void}$$

The args buffer holds one 64-bit slot per wasm parameter. One dispatch shape covers wasm-defined functions, host imports, and indirect calls.

The JitExecEnv struct groups state into three sets.

Group	Contents
Lookup tables	FuncTable, linear memory, globals, indirect-call shadow
Helper addresses	host call, memory growth, profile-saturation helpers
Promotion state	call counters, back-edge counters, OsrEntryTable, OSR buffer

JitExecEnv points to shared tables; **Tier-2** publication updates FuncTable slots or OsrEntryTable slots.

Backup: scope of JitExecEnv writes

Field ownership keeps cross-thread synchronization simple.

- The user thread writes every field except the `0srEntryTable` pointer.
- The worker thread writes only the `0srEntryTable` pointer.
- Publishing a **Tier-2** batch is one atomic store into `FuncTable[i]`.
- Publishing an OSR continuation is one atomic store into `0srEntryTable[id]`.

No locks sit on the hot dispatch path. On the supported host architecture, the dispatch load already provides the ordering the design needs.

Backup: full pseudocode of the back-edge diamond

```
slot = LOAD OsrEntryTable[loopId]
if slot != WAITING:                // sentinel 0
    if slot == COUNTING:           // sentinel 1
        counter = LOAD BackEdgeCounters[loopId]
        if counter < threshold:    // unsigned!
            counter += 1
        STORE BackEdgeCounters[loopId] = counter
        if counter == threshold: notify_osr(loopId)
    else:                          // READY: slot holds an entry pointer
        snapshot wasm locals into per-thread buffer
        return slot(env, locals)    // tail-call the OSR entry
fall through to the back-edge
```


Backup: OSR lineage

HotSpot pattern. A per-method back-edge counter triggers an OSR compile. The OSR compile emits a method version whose entry point is the loop header. Live state transfer uses a stack map. HotSpot has supported OSR for over two decades.

.NET 7 (2022). Added OSR with the same one-shot-caller motivation this thesis targets. The release notes describe methods called once that spend their run-time inside loops.

This thesis. Follows the HotSpot and .NET-7 lineage. Three simplifications:

- Restricted to outermost-loop back-edges.
- Live state is the wasm-level local set, not a stack map.
- Continuation is a fresh wasm function compiled through the LLVM frontend.

WebAssembly's structured control flow makes the simplifications safe.

Related work

Browser-engine Wasm pipelines. V8 pairs Liftoff and TurboFan. JavaScriptCore runs IPInt, BBQ, and OMG. SpiderMonkey runs Baseline and Ion. All three use function-entry promotion. Public documentation reports no OSR between the baseline JIT and the optimizing JIT.

Server-side Wasm pipelines. Wasmtime offers Cranelift (optimizing) and Winch (baseline). The runtime does not currently tier up between them.

OSR in language VMs. HotSpot has supported OSR for over two decades. .NET 7 added OSR in 2022 with the one-shot-caller motivation this thesis also targets. Tracing JITs (PyPy, LuaJIT) sidestep OSR through trace recording.

This thesis fills a gap. No published server-side WebAssembly runtime combines a non-LLVM Sea-of-Nodes baseline, selective LLVM-frontend reuse, and loop-level OSR.

Backup: workload categorization (39 kernels)

Frequent-call (17) blake3-scalar, bz2, gcc-loops, hashset, pulldown-cmark, quicksort, regex, shootout-gimli, shootout-heapsort, shootout-keccak, shootout-memmove, shootout-seqhash, shootout-sieve, shootout-switch, shootout-xblabla20, shootout-xchacha20, tinygo-regex

Hot-leaf (6) blind-sig, rust-compression, rust-json, rust-protobuf, shootout-ctype, shootout-ed25519

One-shot-caller (7) blind-sig, shootout-base64, shootout-ctype, shootout-matrix, shootout-minicsv, shootout-random, shootout-ratelimit

Resistant (4) richards, rust-html-rewriter, shootout-fib2, shootout-nestedloop

Compile-dominated (7) image-classification, spidermonkey-json, spidermonkey-markdown, spidermonkey-regex, sqlite3, tinygo-json, tract-onnx-image-classification

shootout-ctype and blind-sig sit in two groups. Both mechanisms must fire to close the gap.

Backup: bibliography highlights

- Click & Paleczny, *A simple graph-based intermediate representation*, IR '95. – Sea-of-Nodes.
- Click, *Global code motion / global value numbering*, SIGPLAN '95.
- Fink & Qian, *Adaptive recompilation with on-stack replacement*, CGO 2003.
- D'Elia & Demetrescu, *On-stack replacement, distilled*, TOPLAS 2018.
- Bytecode Alliance, *Wasmtime baseline compilation RFC*, 2022.
- Clark, *Firefox's streaming and tiering compiler*, Mozilla Hacks, 2018.
- Chen et al., *Accelerate RISC-V instruction set simulation by tiered JIT compilation*, VMIL 2024.
- Stogov, *dstogov/ir*: Sea-of-Nodes JIT library.